

Plagiarism detection is the process of locating instances of plagiarism within a document or a body of work. Automated similarity detection systems compare the text of submitted documents against one another to find passages that are suspiciously similar. A typical pipeline first extracts the raw text, then normalizes it by lowercasing and removing punctuation, and finally compares the documents using a similarity algorithm such as cosine similarity or the Jaccard index. When the computed score exceeds a configured threshold, the pair of documents is flagged for manual review by an instructor.